

Individual Assignment 2 Task 1

Name: Rohit Panda

UOW ID: 8943060

Before Running the code, ensure you have the kaggle.json file to upload into the Colab.

Download the kaggle.json from this link: [Download Now](#)

```
In [ ]: !pip install -q kaggle
```

```
In [ ]: # 1. Upload kaggle.json
from google.colab import files
files.upload()

# 2. Move it to the correct directory and set permissions
!mkdir -p ~/.kaggle
!mv kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json

!kaggle datasets download -d nikhil7280/weather-type-classification
!unzip -o weather-type-classification.zip

import pandas as pd

df = pd.read_csv('weather_classification_data.csv')
print("Dataset loaded successfully!")
df.head()
```

No file chosen

Upload widget is only available when the cell

has been executed in the current browser session. Please rerun this cell to enable.

Saving kaggle.json to kaggle.json

Dataset URL: <https://www.kaggle.com/datasets/nikhil7280/weather-type-classification>

License(s): other

weather-type-classification.zip: Skipping, found more recently modified local copy (use --force to force download)

Archive: weather-type-classification.zip

inflating: weather_classification_data.csv

Dataset loaded successfully!

Out[]:

	Temperature	Humidity	Wind Speed	Precipitation (%)	Cloud Cover	Atmospheric Pressure	UV Index
0	14.0	73	9.5	82.0	partly cloudy	1010.82	2
1	39.0	96	8.5	71.0	partly cloudy	1011.43	7
2	30.0	64	7.0	16.0	clear	1018.72	5
3	38.0	83	1.5	82.0	clear	1026.25	7
4	27.0	74	17.0	66.0	overcast	990.67	1

In []:

```
# CSCI316 - Task 1: Naïve Bayes Weather Type Classification
# Individual Assignment 2 - 2025 Session 3 (SIM)

## 1. Import Required Libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, StratifiedShuffleSplit
from sklearn.preprocessing import OneHotEncoder, LabelEncoder
from sklearn.naive_bayes import GaussianNB, MultinomialNB, CategoricalNB
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.compose import ColumnTransformer
import warnings
warnings.filterwarnings('ignore')
```

In []:

```
# 2. Load the real dataset from Kaggle
file_name = 'weather_classification_data.csv' # Adjust if the actual file name is different
if os.path.exists(file_name):
    df = pd.read_csv(file_name)
    print("Dataset loaded successfully!")
else:
    raise FileNotFoundError("weather_classification_data.csv not found. Please check the file name and path.")

# Sample weather data
df = pd.DataFrame({
    'Temperature': np.random.normal(25, 10, n_samples),
    'Humidity': np.random.normal(60, 20, n_samples),
    'Wind Speed': np.random.normal(15, 5, n_samples),
    'Precipitation': np.random.uniform(0, 100, n_samples),
    'Cloud Cover': np.random.choice(['Clear', 'Partly Cloudy', 'Overcast'], n_samples),
    'Atmospheric Pressure': np.random.normal(1013, 20, n_samples),
    'UV Index': np.random.uniform(0, 11, n_samples),
    'Season': np.random.choice(['Spring', 'Summer', 'Autumn', 'Winter'], n_samples),
    'Visibility': np.random.uniform(1, 25, n_samples),
    'Location': np.random.choice(['Coastal', 'Inland', 'Mountain'], n_samples),
    'Weather Type': np.random.choice(['Sunny', 'Rainy', 'Cloudy', 'Snowy'], n_samples)
})
print("Sample dataset created for demonstration")
```

```
print(f"Dataset shape: {df.shape}")
print(f"Dataset columns: {df.columns.tolist()}")
```

Dataset loaded successfully!

Dataset shape: (13200, 11)

Dataset columns: ['Temperature', 'Humidity', 'Wind Speed', 'Precipitation (%)', 'Cloud Cover', 'Atmospheric Pressure', 'UV Index', 'Season', 'Visibility (km)', 'Location', 'Weather Type']

In []: *## 3. Data Exploration and Visualization*

```
# Display basic information about the dataset
```

```
print("\n=== Dataset Information ===")
```

```
print(df.info())
```

```
print("\n=== First 5 rows ===")
```

```
print(df.head())
```

```
# Check for missing values
```

```
print("\n=== Missing Values ===")
```

```
print(df.isnull().sum())
```

```
# Statistical summary
```

```
print("\n=== Statistical Summary ===")
```

```
print(df.describe())
```

```
# Target variable distribution
```

```
print("\n=== Weather Type Distribution ===")
```

```
print(df['Weather Type'].value_counts())
```

```
# Visualize target variable distribution
```

```
plt.figure(figsize=(10, 6))
```

```
plt.subplot(1, 2, 1)
```

```
df['Weather Type'].value_counts().plot(kind='bar')
```

```
plt.title('Weather Type Distribution')
```

```
plt.xlabel('Weather Type')
```

```
plt.ylabel('Count')
```

```
plt.xticks(rotation=45)
```

```
plt.subplot(1, 2, 2)
```

```
df['Weather Type'].value_counts().plot(kind='pie', autopct='%1.1f%%')
```

```
plt.title('Weather Type Distribution (Pie Chart)')
```

```
plt.ylabel('')
```

```
plt.tight_layout()
```

```
plt.show()
```

```
# Correlation matrix for numerical features
```

```
numerical_features = df.select_dtypes(include=[np.number]).columns
```

```
if len(numerical_features) > 1:
```

```
    plt.figure(figsize=(10, 8))
```

```
    correlation_matrix = df[numerical_features].corr()
```

```
    sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0)
```

```
    plt.title('Correlation Matrix of Numerical Features')
```

```
    plt.show()
```

=== Dataset Information ===

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 13200 entries, 0 to 13199

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	Temperature	13200 non-null	float64
1	Humidity	13200 non-null	int64
2	Wind Speed	13200 non-null	float64
3	Precipitation (%)	13200 non-null	float64
4	Cloud Cover	13200 non-null	object
5	Atmospheric Pressure	13200 non-null	float64
6	UV Index	13200 non-null	int64
7	Season	13200 non-null	object
8	Visibility (km)	13200 non-null	float64
9	Location	13200 non-null	object
10	Weather Type	13200 non-null	object

dtypes: float64(5), int64(2), object(4)

memory usage: 1.1+ MB

None

=== First 5 rows ===

	Temperature	Humidity	Wind Speed	Precipitation (%)	Cloud Cover \
0	14.0	73	9.5	82.0	partly cloudy
1	39.0	96	8.5	71.0	partly cloudy
2	30.0	64	7.0	16.0	clear
3	38.0	83	1.5	82.0	clear
4	27.0	74	17.0	66.0	overcast

	Atmospheric Pressure	UV Index	Season	Visibility (km)	Location \
0	1010.82	2	Winter	3.5	inland
1	1011.43	7	Spring	10.0	inland
2	1018.72	5	Spring	5.5	mountain
3	1026.25	7	Spring	1.0	coastal
4	990.67	1	Winter	2.5	mountain

	Weather Type
0	Rainy
1	Cloudy
2	Sunny
3	Sunny
4	Rainy

=== Missing Values ===

Temperature	0
Humidity	0
Wind Speed	0
Precipitation (%)	0
Cloud Cover	0
Atmospheric Pressure	0
UV Index	0
Season	0
Visibility (km)	0
Location	0
Weather Type	0

dtype: int64

=== Statistical Summary ===

	Temperature	Humidity	Wind Speed	Precipitation (%) \
count	13200.000000	13200.000000	13200.000000	13200.000000
mean	19.127576	68.710833	9.832197	53.644394
std	17.386327	20.194248	6.908704	31.946541
min	-25.000000	20.000000	0.000000	0.000000
25%	4.000000	57.000000	5.000000	19.000000
50%	21.000000	70.000000	9.000000	58.000000
75%	31.000000	84.000000	13.500000	82.000000
max	109.000000	109.000000	48.500000	109.000000

	Atmospheric Pressure	UV Index	Visibility (km)
count	13200.000000	13200.000000	13200.000000
mean	1005.827896	4.005758	5.462917
std	37.199589	3.856600	3.371499
min	800.120000	0.000000	0.000000
25%	994.800000	1.000000	3.000000
50%	1007.650000	3.000000	5.000000
75%	1016.772500	7.000000	7.500000
max	1199.210000	14.000000	20.000000

=== Weather Type Distribution ===

Weather Type



Rainy 3300


Cloudy 3300

Sunny 3300

Snowy 3300

Name: count, dtype: int64





```
In [ ]: ## 4. Data Preprocessing

# Identify categorical and numerical features
categorical_features = df.select_dtypes(include=['object']).columns.tolist()
numerical_features = df.select_dtypes(include=[np.number]).columns.tolist()

# Remove target variable from features
if 'Weather Type' in categorical_features:
    categorical_features.remove('Weather Type')
if 'Weather Type' in numerical_features:
    numerical_features.remove('Weather Type')

print(f"Categorical features: {categorical_features}")
print(f"Numerical features: {numerical_features}")

# Prepare features and target
X = df.drop('Weather Type', axis=1)
y = df['Weather Type']

# Encode target variable
label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)
```

```
print(f"\nTarget classes: {label_encoder.classes_}")
print(f"Encoded target shape: {y_encoded.shape}")
```

Categorical features: ['Cloud Cover', 'Season', 'Location']

Numerical features: ['Temperature', 'Humidity', 'Wind Speed', 'Precipitation (%)', 'Atmospheric Pressure', 'UV Index', 'Visibility (km)']

Target classes: ['Cloudy' 'Rainy' 'Snowy' 'Sunny']

Encoded target shape: (13200,)

In []: *## 5. Feature Engineering - One-Hot Encoding for Categorical Features*

```
# Create preprocessing pipeline
if categorical_features:
    # One-hot encode categorical features
    preprocessor = ColumnTransformer(
        transformers=[
            ('num', 'passthrough', numerical_features),
            ('cat', OneHotEncoder(drop='first', sparse_output=False), categorical_features)
        ]
    )

    # Fit and transform the features
    X_processed = preprocessor.fit_transform(X)

    # Get feature names after transformation
    feature_names = numerical_features.copy()
    if hasattr(preprocessor.named_transformers_['cat'], 'get_feature_names_out'):
        cat_feature_names = preprocessor.named_transformers_['cat'].get_feature_names_out()
        feature_names.extend(cat_feature_names)

    print(f"Features after preprocessing: {len(feature_names)}")
    print(f"Processed data shape: {X_processed.shape}")
else:
    X_processed = X.values
    feature_names = numerical_features
```

Features after preprocessing: 15

Processed data shape: (13200, 15)

In []: *## 6. Stratified Train-Test Split*

```
# Use stratified sampling to maintain class distribution
stratified_split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
train_idx, test_idx = next(stratified_split.split(X_processed, y_encoded))

X_train = X_processed[train_idx]
X_test = X_processed[test_idx]
y_train = y_encoded[train_idx]
y_test = y_encoded[test_idx]

print(f"\nTraining set shape: {X_train.shape}")
print(f"Test set shape: {X_test.shape}")

# Check class distribution in train and test sets
print(f"\nTraining set class distribution:")
unique_train, counts_train = np.unique(y_train, return_counts=True)
```

```

for i, count in enumerate(counts_train):
    print(f" {label_encoder.classes_[unique_train[i]]}: {count} ({count/len(counts_train)})")

print(f"\nTest set class distribution:")
unique_test, counts_test = np.unique(y_test, return_counts=True)
for i, count in enumerate(counts_test):
    print(f" {label_encoder.classes_[unique_test[i]]}: {count} ({count/len(counts_test)})")

```

Training set shape: (10560, 15)

Test set shape: (2640, 15)

Training set class distribution:

```

Cloudy: 2640 (25.0%)
Rainy: 2640 (25.0%)
Snowy: 2640 (25.0%)
Sunny: 2640 (25.0%)

```

Test set class distribution:

```

Cloudy: 660 (25.0%)
Rainy: 660 (25.0%)
Snowy: 660 (25.0%)
Sunny: 660 (25.0%)

```

In []: *## 7. Naïve Bayes Model Implementation*

```

# Try different Naïve Bayes variants
models = {
    'Gaussian Naïve Bayes': GaussianNB(),
    'Multinomial Naïve Bayes': MultinomialNB(),
}

results = {}

for name, model in models.items():
    print(f"\n=== {name} ===")

    try:
        # Train the model
        model.fit(X_train, y_train)

        # Make predictions
        y_train_pred = model.predict(X_train)
        y_test_pred = model.predict(X_test)

        # Calculate accuracies
        train_accuracy = accuracy_score(y_train, y_train_pred)
        test_accuracy = accuracy_score(y_test, y_test_pred)

        print(f"Training Accuracy: {train_accuracy:.4f}")
        print(f"Test Accuracy: {test_accuracy:.4f}")

        # Store results
        results[name] = {
            'model': model,
            'train_accuracy': train_accuracy,
            'test_accuracy': test_accuracy,

```



```

        'train_predictions': y_train_pred,
        'test_predictions': y_test_pred
    }

```

```

except Exception as e:
    print(f"Error with {name}: {str(e)}")

```

=== Gaussian Naïve Bayes ===

Training Accuracy: 0.8290

Test Accuracy: 0.8235

=== Multinomial Naïve Bayes ===

Error with Multinomial Naïve Bayes: Negative values in data passed to MultinomialNB (input X).

In []: *## 8. Model Evaluation and Results*

```

# Select the best performing model
best_model_name = max(results.keys(), key=lambda x: results[x]['test_accuracy'])
best_model = results[best_model_name]['model']

print(f"\n=== Best Model: {best_model_name} ===")
print(f"Training Accuracy: {results[best_model_name]['train_accuracy']:.4f}")
print(f"Test Accuracy: {results[best_model_name]['test_accuracy']:.4f}")

# Detailed classification report
print(f"\n=== Classification Report for {best_model_name} ===")
y_test_pred_best = results[best_model_name]['test_predictions']
report = classification_report(y_test, y_test_pred_best,
                              target_names=label_encoder.classes_,
                              output_dict=True)

print(classification_report(y_test, y_test_pred_best,
                           target_names=label_encoder.classes_))

# Confusion Matrix
cm = confusion_matrix(y_test, y_test_pred_best)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=label_encoder.classes_,
            yticklabels=label_encoder.classes_)
plt.title(f'Confusion Matrix - {best_model_name}')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```


=== Best Model: Gaussian Naïve Bayes ===

Training Accuracy: 0.8290

Test Accuracy: 0.8235

=== Classification Report for Gaussian Naïve Bayes ===

	precision	recall	f1-score	support
Cloudy	0.88	0.77	0.82	660
Rainy	0.78	0.75	0.76	660
Snowy	0.80	0.90	0.85	660
Sunny	0.84	0.88	0.86	660
accuracy			0.82	2640
macro avg	0.83	0.82	0.82	2640
weighted avg	0.83	0.82	0.82	2640



```
In [ ]: ## 9. Feature Importance Analysis (for Gaussian NB)
```

```
if best_model_name == 'Gaussian Naïve Bayes':  
    # For Gaussian NB, we can analyze feature importance based on variance  
    print(f"\n=== Feature Analysis for Gaussian Naïve Bayes ===")  
  
    # Get class-wise feature statistics  
    n_classes = len(label_encoder.classes_)  
    n_features = X_train.shape[1]
```

```

print(f"Number of classes: {n_classes}")
print(f"Number of features: {n_features}")

# Display theta (mean) and sigma (variance) for each class
if hasattr(best_model, 'theta_') and hasattr(best_model, 'var_'):
    print("\nClass-wise feature means (theta):")
    for i, class_name in enumerate(label_encoder.classes_):
        print(f"\n{class_name}:")
        if len(feature_names) == len(best_model.theta_[i]):
            for j, feature in enumerate(feature_names[:min(10, len(feature_names))]):
                print(f"    {feature}: {best_model.theta_[i][j]:.4f}")
        else:
            print(f"    Mean values: {best_model.theta_[i][:5]}..." # Sh

```

=== Feature Analysis for Gaussian Naïve Bayes ===

Number of classes: 4

Number of features: 15

Class-wise feature means (theta):

Cloudy:

Temperature: 22.7985
Humidity: 66.3155
Wind Speed: 8.6400
Precipitation (%): 40.3705
Atmospheric Pressure: 1010.0687
UV Index: 3.5742
Visibility (km): 7.1225
Cloud Cover_cloudy: 0.0280
Cloud Cover_overcast: 0.3924
Cloud Cover_partly cloudy: 0.5795

Rainy:

Temperature: 22.7595
Humidity: 78.2015
Wind Speed: 13.7523
Precipitation (%): 74.7394
Atmospheric Pressure: 1003.5143
UV Index: 2.6864
Visibility (km): 3.6337
Cloud Cover_cloudy: 0.0307
Cloud Cover_overcast: 0.6583
Cloud Cover_partly cloudy: 0.3110

Snowy:

Temperature: -1.5864
Humidity: 78.5034
Wind Speed: 11.0536
Precipitation (%): 74.7652
Atmospheric Pressure: 991.2356
UV Index: 1.9242
Visibility (km): 3.5602
Cloud Cover_cloudy: 0.0318
Cloud Cover_overcast: 0.7663
Cloud Cover_partly cloudy: 0.2019

Sunny:

Temperature: 32.3511
Humidity: 51.2159
Wind Speed: 6.0852
Precipitation (%): 24.5466
Atmospheric Pressure: 1017.8063
UV Index: 7.8193
Visibility (km): 7.5525
Cloud Cover_cloudy: 0.0330
Cloud Cover_overcast: 0.0299
Cloud Cover_partly cloudy: 0.2917

In []: *## 10. Model Comparison Visualization*

```

if len(results) > 1:
    # Compare model performances
    model_names = list(results.keys())
    train_accuracies = [results[name]['train_accuracy'] for name in model_names]
    test_accuracies = [results[name]['test_accuracy'] for name in model_names]

    x = np.arange(len(model_names))
    width = 0.35

    plt.figure(figsize=(10, 6))
    plt.bar(x - width/2, train_accuracies, width, label='Training Accuracy', color='green')
    plt.bar(x + width/2, test_accuracies, width, label='Test Accuracy', color='red')

    plt.xlabel('Models')
    plt.ylabel('Accuracy')
    plt.title('Model Performance Comparison')
    plt.xticks(x, model_names, rotation=45)
    plt.legend()
    plt.tight_layout()
    plt.show()

```

In []: *## 11. Conclusions and Summary*

```

print("\n=== SUMMARY AND CONCLUSIONS ===")
print(f"1. Dataset contains {df.shape[0]} samples with {df.shape[1]} features")
print(f"2. Target variable has {len(label_encoder.classes_)} classes: {label_encoder.classes_}")
print(f"3. Used stratified sampling: {len(X_train)} training samples, {len(X_test)} test samples")
print(f"4. Applied one-hot encoding to categorical features")
print(f"5. Best performing model: {best_model_name}")
print(f"6. Final test accuracy: {results[best_model_name]['test_accuracy']:.4f}")

# Performance metrics summary
print(f"\n=== PERFORMANCE METRICS ===")
for metric in ['precision', 'recall', 'f1-score']:
    macro_avg = report['macro avg'][metric]
    weighted_avg = report['weighted avg'][metric]
    print(f"{metric.title()}: Macro avg = {macro_avg:.4f}, Weighted avg = {weighted_avg:.4f}")

print(f"\n=== IMPLEMENTATION NOTES ===")
print("1. Stratified sampling ensures balanced class distribution in train/test sets")
print("2. One-hot encoding transforms categorical features to numerical form")
print("3. Gaussian Naïve Bayes works well with continuous features")
print("4. Model assumes feature independence (Naïve Bayes assumption)")
print("5. Performance can be improved with feature selection and hyperparameter tuning")

print("\n=== TASK 1 COMPLETED SUCCESSFULLY ===")

```

=== SUMMARY AND CONCLUSIONS ===

1. Dataset contains 13200 samples with 11 features
2. Target variable has 4 classes: Cloudy, Rainy, Snowy, Sunny
3. Used stratified sampling: 10560 training samples, 2640 test samples
4. Applied one-hot encoding to categorical features
5. Best performing model: Gaussian Naïve Bayes
6. Final test accuracy: 0.8235

=== PERFORMANCE METRICS ===

Precision: Macro avg = 0.8256, Weighted avg = 0.8256

Recall: Macro avg = 0.8235, Weighted avg = 0.8235

F1-Score: Macro avg = 0.8226, Weighted avg = 0.8226

=== IMPLEMENTATION NOTES ===

1. Stratified sampling ensures balanced class distribution in train/test splits
2. One-hot encoding transforms categorical features to numerical format
3. Gaussian Naïve Bayes works well with continuous features
4. Model assumes feature independence (Naïve Bayes assumption)
5. Performance can be improved with feature selection and hyperparameter tuning

=== TASK 1 COMPLETED SUCCESSFULLY ===

This notebook was converted with convert.ploomber.io