

# Individual Assignment 2 Task 2

Name: Rohit Panda

UOW ID: 8943060

Before Running the code, ensure you have the kaggle.json file to upload into the Colab.

Download the kaggle.json from this link: [Download Now](#)

```
In [ ]: !pip install -q kaggle
```

```
In [ ]: # 1. Upload kaggle.json
from google.colab import files
files.upload()

# 2. Move it to the correct directory and set permissions
!mkdir -p ~/.kaggle
!mv kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json

!kaggle datasets download -d nikhil7280/weather-type-classification
!unzip -o weather-type-classification.zip

import pandas as pd

df = pd.read_csv('weather_classification_data.csv')
print("Dataset loaded successfully!")
df.head()
```

No file chosen

Upload widget is only available when the cell

has been executed in the current browser session. Please rerun this cell to enable.

Saving kaggle.json to kaggle.json

Dataset URL: <https://www.kaggle.com/datasets/nikhil7280/weather-type-classification>

License(s): other

Downloading weather-type-classification.zip to /content

0% 0.00/186k [00:00<?, ?B/s]

100% 186k/186k [00:00<00:00, 329MB/s]

Archive: weather-type-classification.zip

inflating: weather\_classification\_data.csv

Dataset loaded successfully!

Out[ ]:

	Temperature	Humidity	Wind Speed	Precipitation (%)	Cloud Cover	Atmospheric Pressure	UV Index
0	14.0	73	9.5	82.0	partly cloudy	1010.82	2
1	39.0	96	8.5	71.0	partly cloudy	1011.43	7
2	30.0	64	7.0	16.0	clear	1018.72	5
3	38.0	83	1.5	82.0	clear	1026.25	7
4	27.0	74	17.0	66.0	overcast	990.67	1

In [ ]:

```
# CSCI316 - Task 2: Artificial Neural Network Weather Type Classification
# Individual Assignment 2 - 2025 Session 3 (SIM)

## 1. Import Required Libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import StratifiedShuffleSplit
from sklearn.preprocessing import StandardScaler, LabelEncoder, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.metrics import classification_report, confusion_matrix
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras.utils import to_categorical
import warnings
warnings.filterwarnings('ignore')

np.random.seed(42)
tf.random.set_seed(42)

print(f"TensorFlow version: {tf.__version__}")
```

TensorFlow version: 2.18.0

In [ ]:

```
## 2. Load Dataset (you may replace with kaggle loading)

try:
    df = pd.read_csv('weather_classification_data.csv')
    print("Dataset loaded successfully!")
except:
    print("Please ensure the dataset is uploaded.")
    n_samples = 1000
    df = pd.DataFrame({
        'Temperature': np.random.normal(25, 10, n_samples),
        'Humidity': np.random.normal(60, 20, n_samples),
        'Wind Speed': np.random.normal(15, 5, n_samples),
        'Precipitation': np.random.uniform(0, 100, n_samples),
        'Cloud Cover': np.random.choice(['Clear', 'Partly Cloudy', 'Overcast', 'Rainy', 'Snowy'], n_samples)
```

```

    'Atmospheric Pressure': np.random.normal(1013, 20, n_samples),
    'UV Index': np.random.uniform(0, 11, n_samples),
    'Season': np.random.choice(['Spring', 'Summer', 'Autumn', 'Winter'],
    'Visibility': np.random.uniform(1, 25, n_samples),
    'Location': np.random.choice(['Coastal', 'Inland', 'Mountain'], n_sa
    'Weather Type': np.random.choice(['Sunny', 'Rainy', 'Cloudy', 'Snowy
    })

```

Dataset loaded successfully!

In [ ]: *## 3. Preprocessing*

```

categorical_features = df.select_dtypes(include='object').columns.tolist()
numerical_features = df.select_dtypes(include=np.number).columns.tolist()
categorical_features.remove('Weather Type') if 'Weather Type' in categorical

X = df.drop('Weather Type', axis=1)
y = df['Weather Type']

label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)
y_categorical = to_categorical(y_encoded)
n_classes = y_categorical.shape[1]

preprocessor = ColumnTransformer([
    ('num', StandardScaler(), numerical_features),
    ('cat', OneHotEncoder(drop='first', sparse_output=False), categorical_fe
])

X_processed = preprocessor.fit_transform(X)

```

In [ ]: *## 4. Train-Test Split*

```

splitter = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
train_idx, test_idx = next(splitter.split(X_processed, y_encoded))

X_train, X_test = X_processed[train_idx], X_processed[test_idx]
y_train, y_test = y_categorical[train_idx], y_categorical[test_idx]
y_train_labels, y_test_labels = y_encoded[train_idx], y_encoded[test_idx]

```

In [ ]: *## 5. ANN Definition*

```

def create_ann_model(input_dim, n_classes, n_hidden_layers=2, n_neurons=64,
    model = keras.Sequential()
    model.add(layers.Dense(n_neurons, input_dim=input_dim, activation='relu',
                           kernel_regularizer=keras.regularizers.l1_l2(l1=1e-4, l2=1e-4)))
    model.add(layers.Dropout(0.3))
    for _ in range(n_hidden_layers - 1):
        model.add(layers.Dense(n_neurons, activation='relu',
                               kernel_regularizer=keras.regularizers.l1_l2(l1=1e-4, l2=1e-4)))
        model.add(layers.Dropout(0.3))
    model.add(layers.Dense(n_classes, activation='softmax'))
    return model

```

In [ ]: *## 6. Grid Search*

```

param_grid = {
    'n_hidden_layers': [2, 3],
    'n_neurons': [64, 128],
    'l1_reg': [0.001, 0.01],
    'l2_reg': [0.001, 0.01]
}

best_score = 0
best_model = None
best_params = {}
grid_results = []

for h in param_grid['n_hidden_layers']:
    for n in param_grid['n_neurons']:
        for l1 in param_grid['l1_reg']:
            for l2 in param_grid['l2_reg']:
                print(f"Testing: layers={h}, neurons={n}, L1={l1}, L2={l2}")
                model = create_ann_model(X_train.shape[1], n_classes, h, n,
                                         model.compile(optimizer='adam', loss='categorical_crossentropy',
                                                         history = model.fit(X_train, y_train, validation_split=0.2,
                                                         epochs=50, batch_size=32, verbose=0,
                                                         callbacks=[EarlyStopping(patience=10, reduce_lr_on_plateau=ReduceLROnPlateau(patience=5,
                                                         loss, acc = model.evaluate(X_test, y_test, verbose=0)
                print(f"Test Accuracy: {acc:.4f}")
                grid_results.append({
                    'n_hidden_layers': h, 'n_neurons': n, 'l1_reg': l1, 'l2_reg': l2,
                    'test_accuracy': acc, 'test_loss': loss
                })
                if acc > best_score:
                    best_score = acc
                    best_params = {'n_hidden_layers': h, 'n_neurons': n, 'l1_reg': l1, 'l2_reg': l2}
                    best_model = model

```

```

Testing: layers=2, neurons=64, L1=0.001, L2=0.001
Test Accuracy: 0.9049
Testing: layers=2, neurons=64, L1=0.001, L2=0.01
Test Accuracy: 0.9015
Testing: layers=2, neurons=64, L1=0.01, L2=0.001
Test Accuracy: 0.8985
Testing: layers=2, neurons=64, L1=0.01, L2=0.01
Test Accuracy: 0.8883
Testing: layers=2, neurons=128, L1=0.001, L2=0.001
Test Accuracy: 0.9023
Testing: layers=2, neurons=128, L1=0.001, L2=0.01
Test Accuracy: 0.9068
Testing: layers=2, neurons=128, L1=0.01, L2=0.001
Test Accuracy: 0.8955
Testing: layers=2, neurons=128, L1=0.01, L2=0.01
Test Accuracy: 0.8977
Testing: layers=3, neurons=64, L1=0.001, L2=0.001
Test Accuracy: 0.9049
Testing: layers=3, neurons=64, L1=0.001, L2=0.01
Test Accuracy: 0.8977
Testing: layers=3, neurons=64, L1=0.01, L2=0.001
Test Accuracy: 0.8970
Testing: layers=3, neurons=64, L1=0.01, L2=0.01
Test Accuracy: 0.8765
Testing: layers=3, neurons=128, L1=0.001, L2=0.001
Test Accuracy: 0.9053
Testing: layers=3, neurons=128, L1=0.001, L2=0.01
Test Accuracy: 0.9004
Testing: layers=3, neurons=128, L1=0.01, L2=0.001
Test Accuracy: 0.8947
Testing: layers=3, neurons=128, L1=0.01, L2=0.01
Test Accuracy: 0.6966

```

In [ ]: *## 7. Final Evaluation*

```

train_loss, train_acc = best_model.evaluate(X_train, y_train, verbose=0)
test_loss, test_acc = best_model.evaluate(X_test, y_test, verbose=0)

y_pred = best_model.predict(X_test)
y_classes = np.argmax(y_pred, axis=1)

print("\n=== FINAL RESULTS ===")
print("Train Accuracy:", train_acc)
print("Test Accuracy:", test_acc)
print("\nBest Hyperparameters:", best_params)
print(classification_report(y_test_labels, y_classes, target_names=label_encoder.classes_)

# Confusion matrix
plt.figure(figsize=(8,6))
sns.heatmap(confusion_matrix(y_test_labels, y_classes), annot=True, fmt='d',
            xticklabels=label_encoder.classes_, yticklabels=label_encoder.classes_)
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()

```

=== FINAL RESULTS ===

Train Accuracy: 0.909469723701477

Test Accuracy: 0.9068182110786438

Best Hyperparameters: {'n\_hidden\_layers': 2, 'n\_neurons': 128, 'l1\_reg': 0.001, 'l2\_reg': 0.01}

	precision	recall	f1-score	support
Cloudy	0.87	0.89	0.88	660
Rainy	0.93	0.90	0.91	660
Snowy	0.98	0.90	0.94	660
Sunny	0.86	0.94	0.90	660
accuracy			0.91	2640
macro avg	0.91	0.91	0.91	2640
weighted avg	0.91	0.91	0.91	2640

