

U

O

W

Classification by Splitting Data

– Dive Into ML Model Training

CSCI316: Big Data Mining Techniques and Implementation



UNIVERSITY
OF WOLLONGONG
AUSTRALIA

Open the black box of model training...

- Recall the following fragment of the end-to-end project (see page 32 of the “End-to-End Big Data Lifecycle” lecture note:

- Try Decision Tree

```
> from sklearn.tree import DecisionTreeRegressor  
> tree_reg = DecisionTreeRegressor()  
> tree_reg.fit(housing_prepared, housing_labels)  
> housing_predictions = tree_reg.predict(housing_prepared)
```

- What is a DT? How does it work? What is the theory behind?*


The Classification Problem: An Example

Tid	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set

Tid	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

Test Set



What is a Decision Tree

- A decision tree is a *flowchart-like tree structure*
 - Each *internal node* (non-leaf node) denotes a test on an attribute
 - Each *branch* (i.e., subtree) represents an outcome of the test
 - Each *leaf node* (or terminal node) holds a class label
- It simulates the process of human decision-making.
 - Thus, one advantage of decision trees is *understandability*

Example of a Decision Tree

<i>Tid</i>	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

categorical

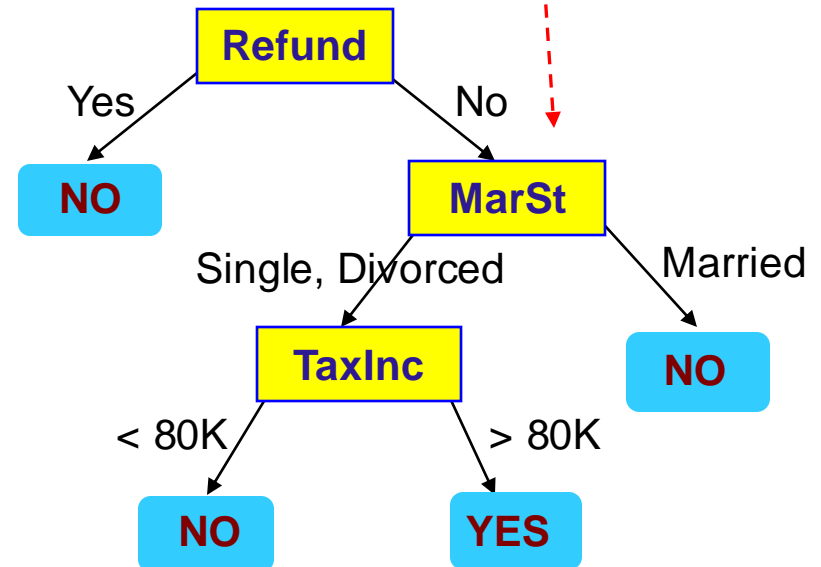
categorical

continuous

class

Each node is associated with a (sub)set of records

Splitting Attributes



Training Data

Model: Decision Tree

Another Example of Decision Tree

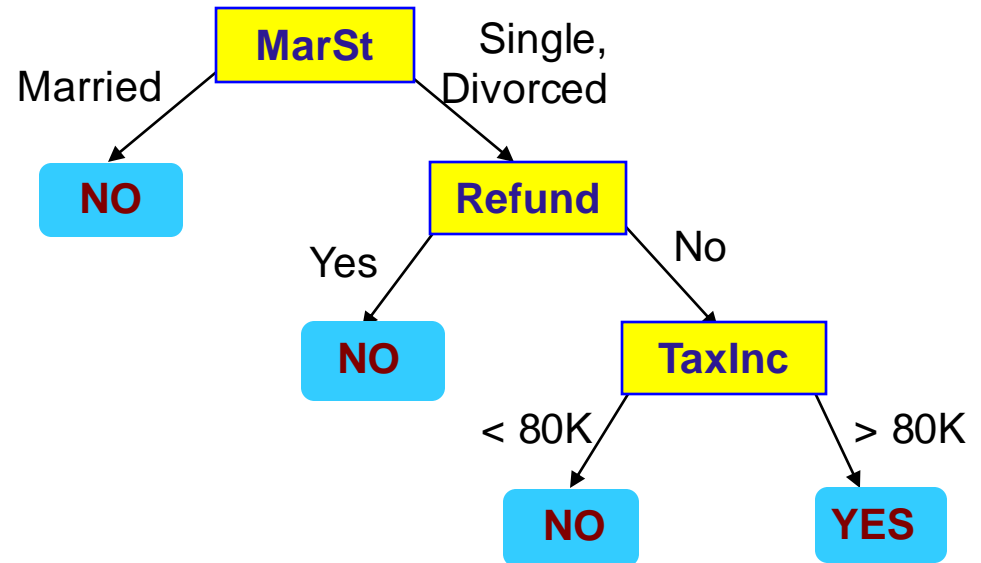
<i>Tid</i>	<i>Refund</i>	<i>Marital Status</i>	<i>Taxable Income</i>	<i>Cheat</i>
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

categorical

categorical

continuous

class



There could be multiple trees that fit the same data!

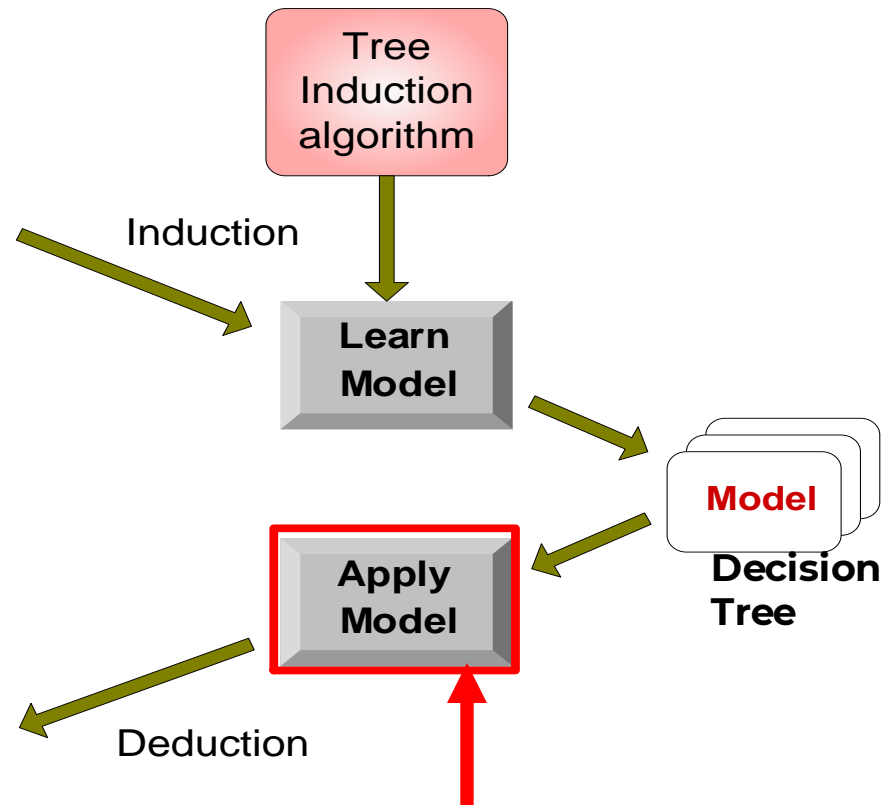
Decision Tree Classification Task

Tid	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set

Tid	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

Test Set

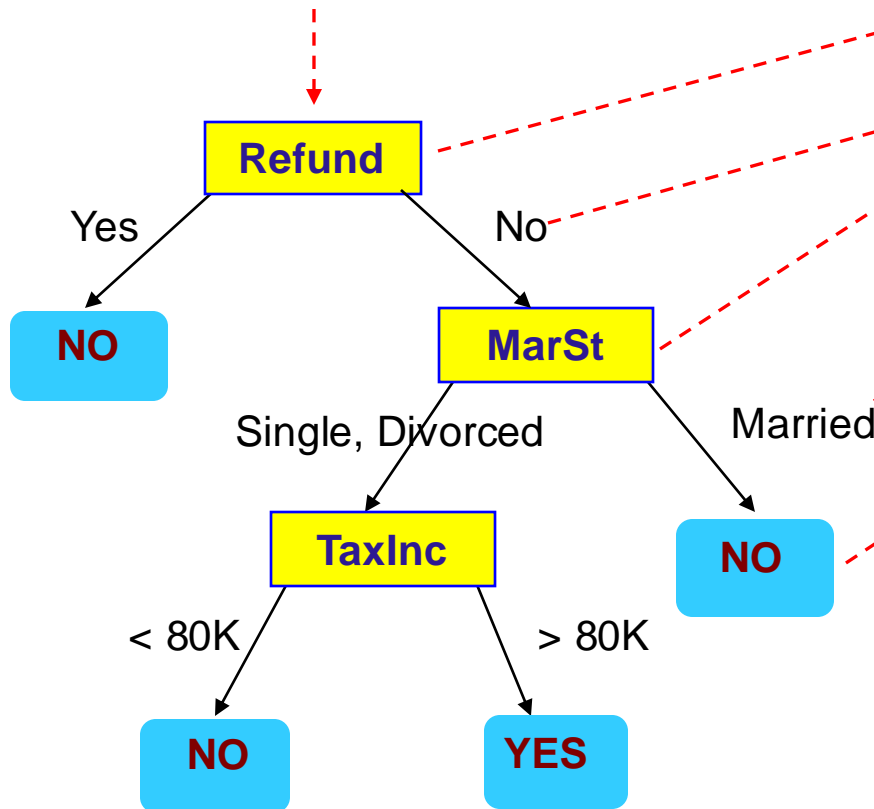


Apply Model to Test Data

Start from the root of tree.

Test Data

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	? No



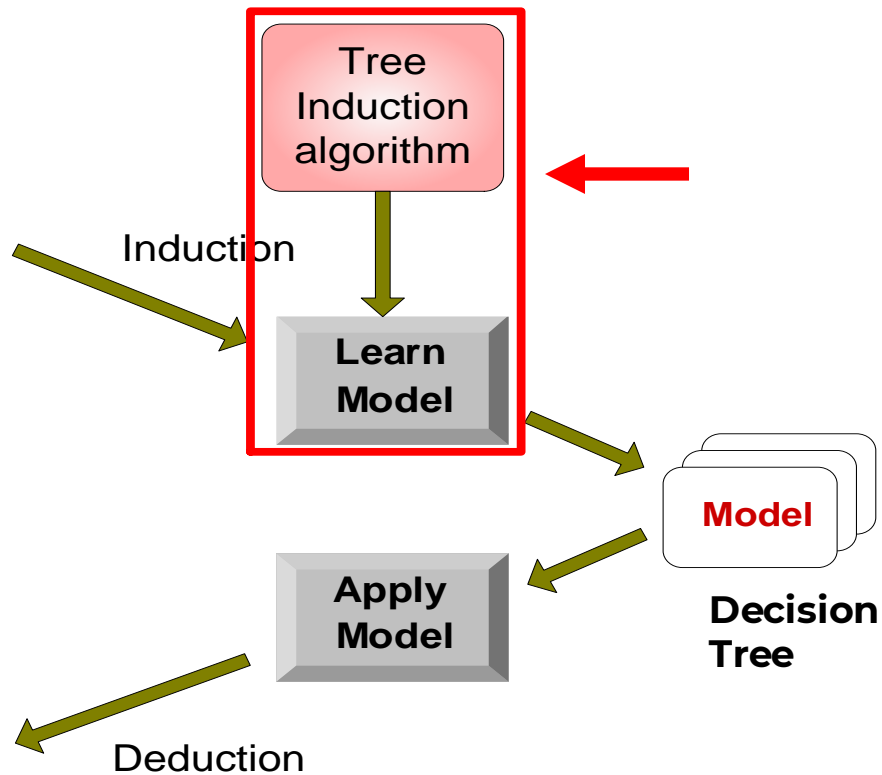
Decision Tree Classification Task

Tid	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set

Tid	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?


Test Set



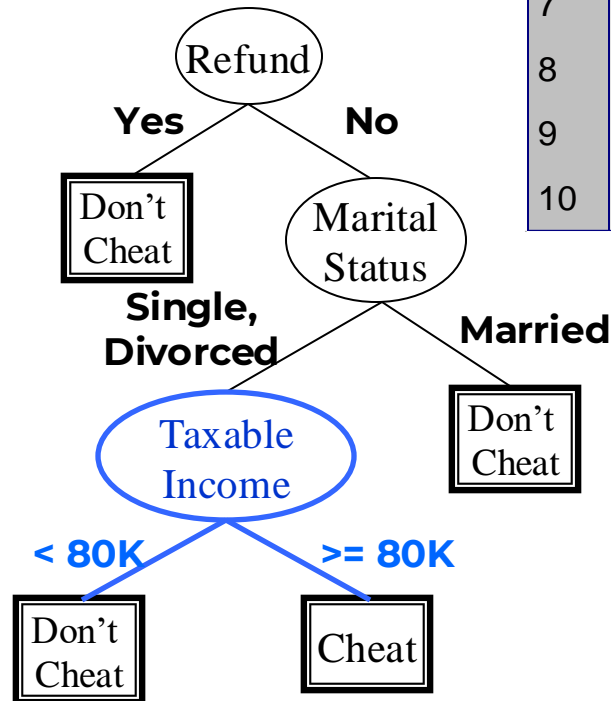
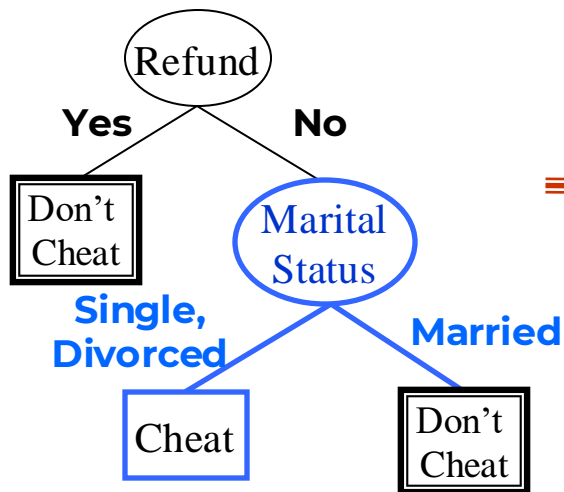

General Structure of Decision Tree Induction Algorithms

- Let D_t be the associated set of training records that reach a node t
- General Procedure:
 - If D_t contains records that belong the same class y_t , then t is a leaf node, labeled as y_t
 - If D_t is an empty set, then t is a leaf node, labeled as the same class as its parent node
 - If no more attributes to split D_t , then t is a leaf node, labeled as the *majority class*
 - Otherwise, **split** the dataset into smaller subsets, each of which is associated with a child node of the node t , and **recursively** apply the same procedure to child node

<i>Tid</i>	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes



Hunt's Algorithm



Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Tree Induction


- Greedy strategy.
 - Split the records based on an attribute test that optimizes certain criterion.
- Issues
 - Determine how to split the records
 - How to specify the attribute test condition? (focus)
 - How to determine the best split?
 - Determine when to stop splitting

How to Specify Test Condition?

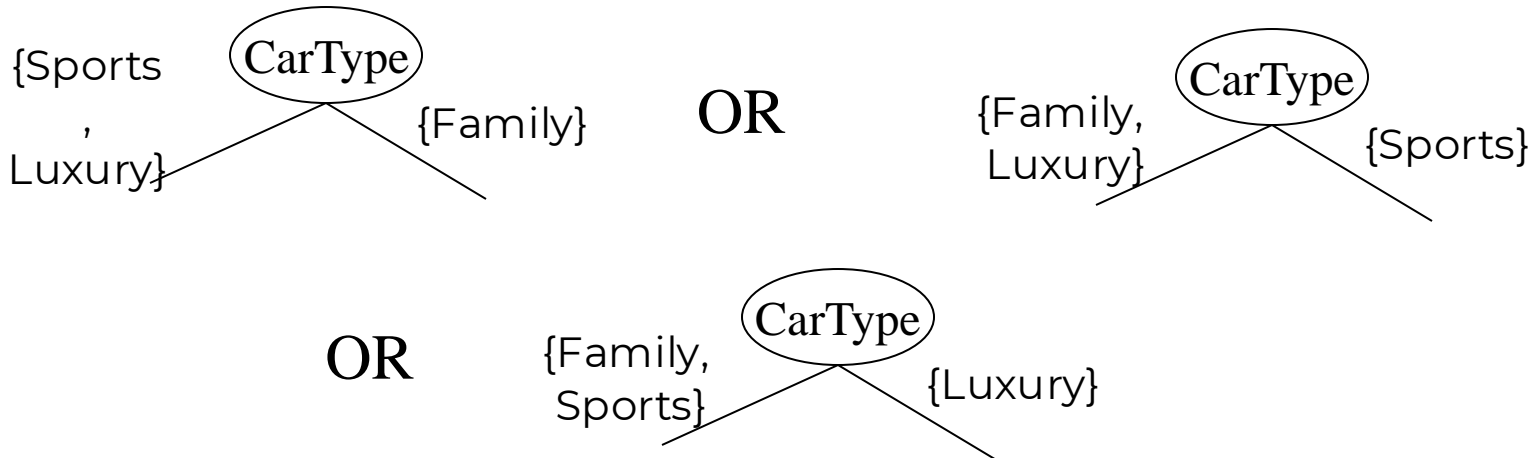
- Depends on the attribute types
 - Nominal/categorical
 - Ordinal
 - Continuous
- Depends on the number of ways to split
 - 2-way split
 - Multi-way split

Splitting Based on Nominal Attributes

- **Multi-way split:** Use as many partitions as distinct values.




- **Binary split:** Divide values into two subsets.
Need to find optimal partitioning.




Splitting Based on Ordinal/Continuous Attributes

- Different ways of handling
 - **Discretization** to form an ordinal categorical attribute
 - Static – discretize once at the beginning
 - Dynamic – bucketing, percentiles, clustering...
 - **Binary Decision**: $(A < v)$ or $(A \geq v)$
 - consider all possible splits and finds the best cut
 - can be more computationally intensive

Splitting Based on Ordinal/Continuous Attributes



(i) Binary split



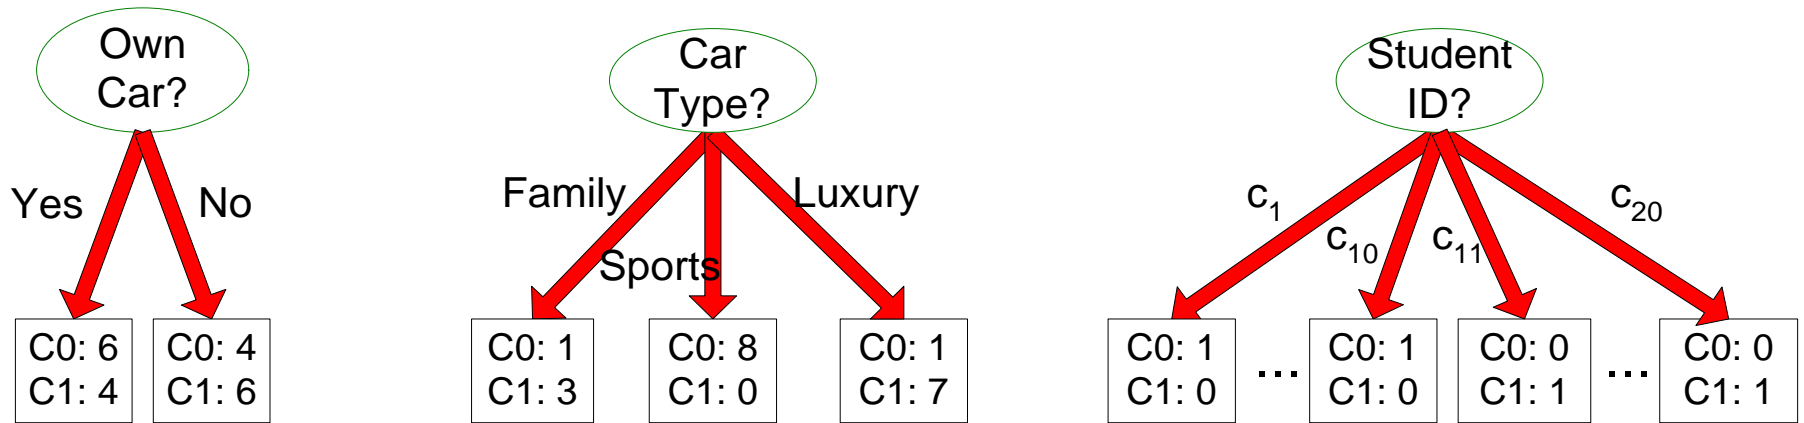
(ii) Multi-way split

Tree Induction

- Greedy strategy.
 - Split the records based on an attribute test that optimizes certain criterion.
- Issues
 - Determine how to split the records
 - How to specify the attribute test condition?
 - How to determine the best split? (focus)
 - Determine when to stop splitting

How to determine the Best Split

**Before Splitting: 10 records of class 0,
10 records of class 1**



Which test condition is the best?

How to determine the Best Split

- Greedy approach:
 - Nodes with **homogeneous** class distributions are preferred
- Need a measure of node **impurity** (or information **uncertainty**):

C0: 5
C1: 5

**Non-homogeneous,
High degree of
impurity**

C0: 9
C1: 1

**Homogeneous,
Low degree of
impurity**

Another way to look at Impurity and Uncertainty

- We flip two different coins: (0 is “head”, 1 is “tail”)
 - 0 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 0 0...
 - 0 1 0 1 0 1 1 1 0 0 1 1 0 1 0 1 0 1...



- Question: *How to measure/quantify the information uncertainty with the two coins?*

Different Measures of Impurity/Uncertainty


- Entropy (information gain)
- Gain ratio
- Gini Index
- Variance
- Others ...

Shannon Entropy

- Logarithm: $y = \log_a x$
 - $2^3 = 8 \Leftrightarrow \log_2 8 = 3$
 - $2^{-1} = 0.5 \Leftrightarrow \log_2 0.5 = -1$
- Shannon Entropy:

$$H(X) = - \sum_{x \in X} p(x) \log_2 p(x)$$

Entropy of a coin:



Conditional Entropy

- Example: $X = \{\text{Raining, Not raining}\}$, $Y = \{\text{Cloudy, not cloudy}\}$

	Cloudy	Not cloudy	Total
Is Raining	24	1	25
Not Raining	25	50	75
Total	49	51	100

- What is the entropy of cloudiness, given the knowledge of whether or not it is raining?

Note. $H(Y|X) \neq H(Y)$

$$\begin{aligned}
 H(Y|X) &= \sum_{x \in X} p(x) H(Y|X = x) \\
 &= \frac{1}{4} H(Y | \text{is raining}) + \frac{3}{4} H(Y | \text{not raining}) \\
 &\approx 0.75 \text{ bits}
 \end{aligned}$$

Information Gain

- If I don't know whether it is raining or not, the entropy of cloudiness is $H(Y) \approx 1.00$ bit (*verifying this as an exercise*)
- How much information about cloudiness do we gain by discovering whether it is raining?
- The Shannon entropy tells $\text{InfoGain}(Y|X) = H(Y) - H(Y|X) \approx 0.25$ bit
- How do we make use of this measure to construct our decision tree?
 - E.g., to determine the **best split** of the dataset.

Splitting Based on InfoGain

- Let D be the set of training records that reach a node
 - Compute the entropy $H(D)$ for D
- Let *Attribute_List* be a set of attributes associated with D
 - Each split with an attribute in *Attribute_List* produces a **partition** on $P = \{D_1, \dots, D_v\}$ on D
 - Compute the conditional entropy for each split and then calculate the InfoGain:

$$H_P(D) = \sum_{i=1}^v \frac{|D_i|}{|D|} H(D_i)$$
$$\text{InfoGain}(P) = H(D) - H_P(D)$$

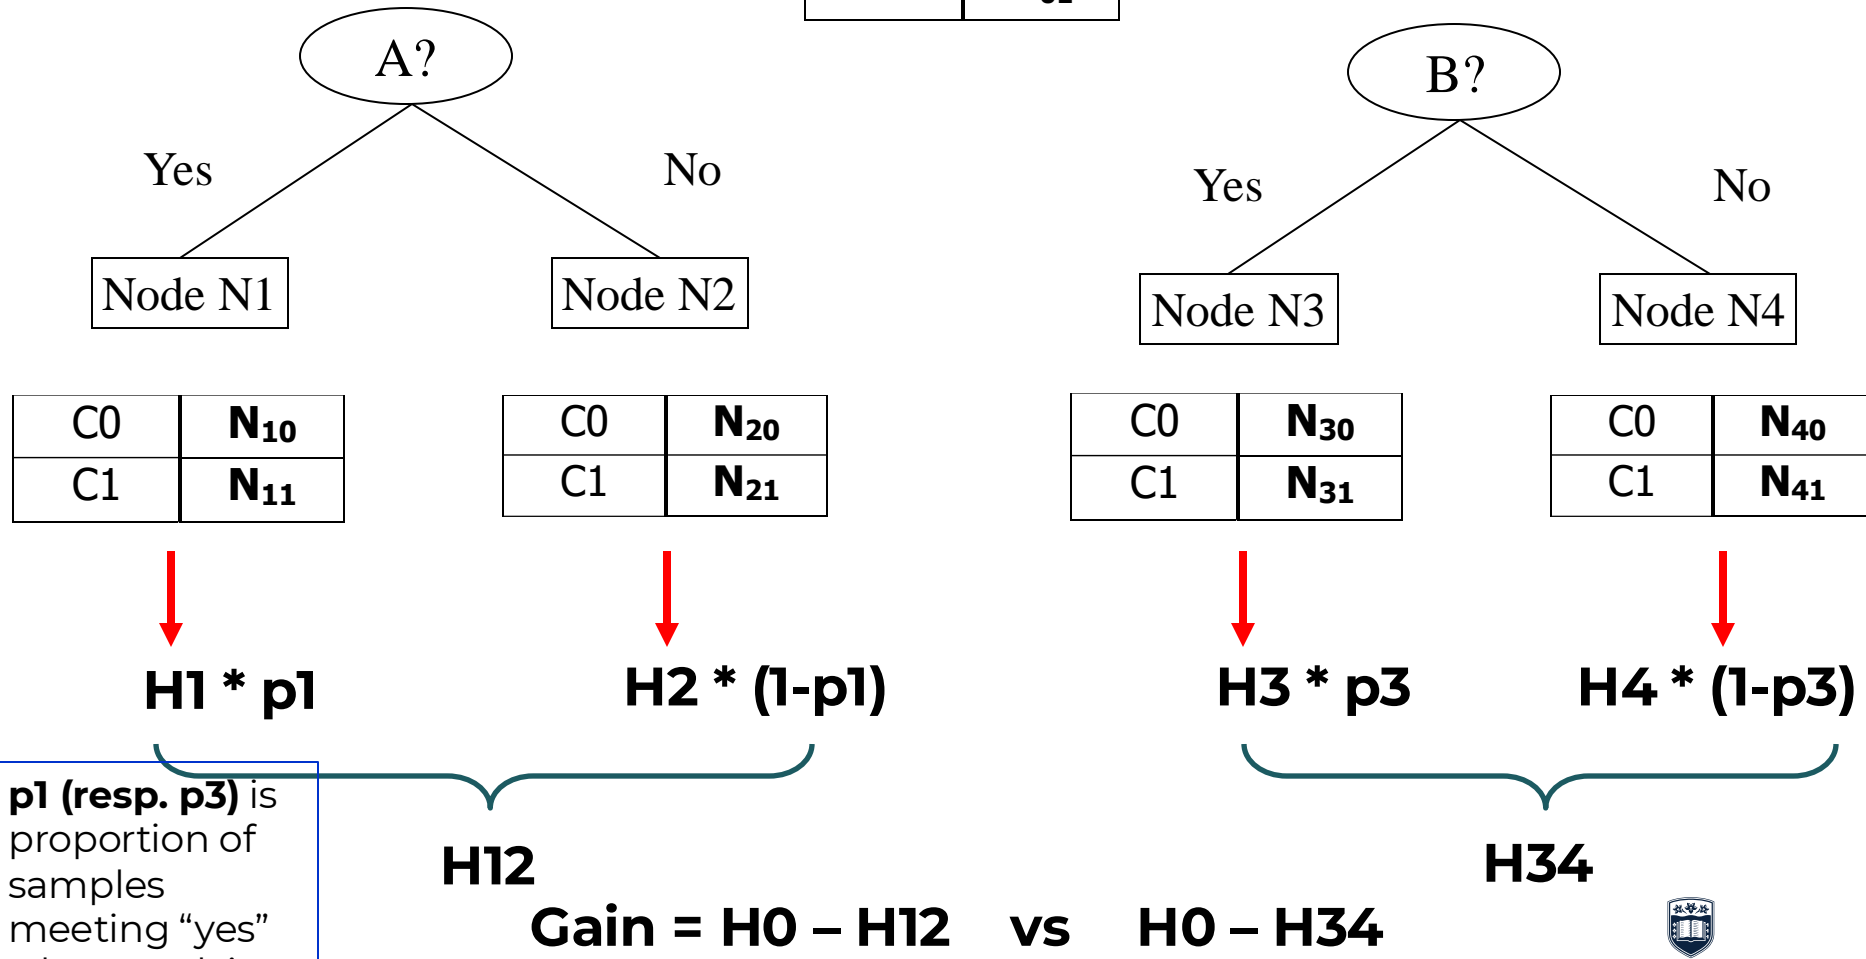
- Select an attribute that gives the best split (one with the *largest* InfoGain)

How to Find the Best Split

Before Splitting:

class	counts
C0	N_{00}
C1	N_{01}

\rightarrow **H0**




$p1$ (resp. $p3$) is proportion of samples meeting "yes" when applying A (resp. B)

Tree Induction

- Greedy strategy.
 - Split the records based on an attribute test that optimizes certain criterion.
- Issues
 - Determine how to split the records
 - How to specify the attribute test condition?
 - How to determine the best split?
 - Determine when to stop splitting (focus)

Stopping Criteria

1. No more attribute for splitting the dataset D_t
 - Majority vote: select the class label with most records to report
2. All tuples in D_t share the same class label
3. D_t is empty (no tuples)
4. Non-basic criteria
 - Tree pre-pruning (talked later), such as
 - set a threshold for the impurity measured
 - minimum dataset size
 - largest tree depth
 - etc.



Tree Induction Algorithm

Assumption: the training tuples contain categorical values only; multi-split is used.

Procedure: **generate_decision_tree**(D , $Attribute_List$).

- ❖ Generate a decision tree from a set of training tuples of D .

Input:

- Dataset, D , which is a set of training tuples (each includes a tuple of feature values and one class label)
- $Attribute_List$, the set of candidate attributes for split

Output: A decision tree

Tree Induction Algorithm

Pseudo-code:

- (1) create a node N ;
- (2) **if** tuples in D are all of the same class, i.e. C , **then**
- (3) **return** N as a (leaf) node labeled with the class C ;
- (4) **if** $Attribute_List$ is empty **then**
- (5) **return** N as a leaf node labeled with the majority class C_0
- (6) **find the *best_splitting_attribute* in $Attribute_List$ to split D ;**
- (7) $New_Attribute_List \leftarrow Attribute_List / \{best_splitting_attribute\}$;

Tree Induction Algorithm

- (8) **foreach** value s of *best_splitting_attribute*;
- (9) let D_s be a subset of D with *best_splitting_attribute* being s ;
- (10) **if** D_s is empty **then**
- (11) attach a (leaf) node labeled with **the majority class in D**
 to node N ;
- (12) **else** attach a new node, N_{child} , returned by applying
 generate_decision_tree(D_s , *New_Attribute_List*) to node N ;
- (13) **return** N ;

Classification with Decision Trees

- Given a testing tuple, the classification with a decision tree is just by traversing the tree until a leaf is reached.
- Procedure: **classify**(N, d)
- Input: testing tuple d .
- Output: a class label C
- Pseudo-code:
 - (1) **if** N is a leaf node **then**
 - (2) **return** the class label C with N ;
 - (3) **else** traverse to the child node N_{child} of N where the value of the *best_splitting_feature* matches the value in d ;
 - (4) let $C = \text{classify}(N_{\text{child}}, d)$;
 - (5) **return** C ;

Python Implementation

- Python **dictionaries** are a convenient data structure to represent a decision tree
 - Each splitting feature is a node
 - For a multi-split tree with categorical features (JSON style):

```
tree = {  
    index_of_splitting_feature: {  
        v_0: subtree_0 or leaf_0,  
        ...  
        v_l: subtree_l or leaf_l  
    }  
}
```

where each v is a (unique) value of the splitting feature.

- Access to the split feature and values:

```
split_feature = tree.keys()  
subtree = tree[split_feature]  
feature_values = subtree.keys()
```

Python Implementation

- A **leaf** can just be a **class label**, say, C_i .
- But more generally, a **leaf** can be represented by a **NumPy array** (i.e., vector) $ary = (q_1, \dots, q_m)$
 - such that $q_i = |D_{C_i}|$ is a **class frequency** where:
 - D is the set of training tuples associated with `splitting_feature` (as a node), and
 - $D_{C_i} \subseteq D$ contains all tuples in D that belong to class C_i
 - Note that a class label can be determined immediately from the vector ary .
 - E.g., just choose the class with the largest q_i

Python Implementation

- It is not hard to observe that both the tree induction and the classification involve a *recursive function*.
- Recursive function example in Python:

```
def factorial(n):  
    if n == 1:  
        return 1  
    else:  
        return n * factorial(n-1)
```

- factorial is called within itself.
- Running:
 $4! = 4 * 3!$
 $3! = 3 * 2!$
 $2! = 2 * 1!$
 $1! = 1$

Python Implementation

- To check whether a node in a tree (as a Python dictionary) is a leaf or grows a subtree:

python3

`isinstance(somenode, dict) == True` *#a subtree*

or

`type(somenode).__name__ == 'dict'` *#a subtree*

Sample Python Code (Compute Shannon Entropy)

```
# calculate Shannon Entropy of a dataset
def calcShannonEnt(dataSet):
    numEntries = len(dataSet) # number of tuples
    labelCounts = {}
    for featVec in dataSet:
        # a class label is the last element in each tuples
        currentLabel = featVec[-1]
        if currentLabel not in labelCounts.keys():
            labelCounts[currentLabel] = 0
        labelCounts[currentLabel] += 1
    shannonEnt = 0.0
    for key in labelCounts:
        prob = float(labelCounts[key]) / numEntries
        shannonEnt -= prob * log(prob, 2)
    return shannonEnt
```

Sample Python Code (Multi-Split , Categorical Features)

```
def chooseBestMultiSplit(dataSet):  
    numFeatures = len(dataSet[0]) - 1 # number of features  
    baseEntropy = calcShannonEnt(dataSet)  
    bestInfoGain = 0.0; bestFeature = -1  
    for i in range(numFeatures): # iterate over all features  
        uniqueVals = set([tuple[i] for tuple in dataSet])  
        newEntropy = 0.0  
        for value in uniqueVals:  
            # "splitDataSet" function, implemented elsewhere, filters "dataset" such that  
            the i-th feature equals to "value"  
            subDataSet = splitDataSet(dataSet, i, value)  
            prob = len(subDataSet) / float(len(dataSet))  
            newEntropy += prob * calcShannonEnt(subDataSet)  
            infoGain = baseEntropy - newEntropy  
            if (infoGain > bestInfoGain):  
                bestInfoGain = infoGain; bestFeature = i  
    return bestFeature # returns a feature index
```

How to Implement a Decision Tree Classifier

- How to represent/encode your decision tree?
 - Consider a Python dictionary (see previous slides)
- How to implement your tree induction algorithm based on the `calcShannonEnt` and `chooseBestMultiSplit` functions?
 - Consider a recursive Python function that calls the two functions
 - Address all basic stopping criteria
- How to classify (new) records with your decision tree?
 - Also consider a recursive function
- *The implementation assumes categorical features, how about ordinal and continuous features?*
 - Use **binning** to generate a suitable number of bins (e.g., 5)

Gini Index

- Gini index (or Gini impurity) is a measure of how often a randomly chosen element from the set would be incorrectly labelled, if it was randomly labelled according to the distribution of labels in the subset.
 - Given D , a set of training tuples:

$$\text{Gini}(D) = \sum_{i=1}^m p_i \sum_{j \neq i} p_j = 1 - \sum_{i=1}^m p_i^2$$

where $p_i = |D_{C_i}|/|D|$, i.e. the probability that a tuple in D belongs to class C_i . (Here D_{C_i} refers to a subset of D such that the tuple belongs to class C_i .)

Gini Index

- For multi-way split on some feature $P = \{D_1, \dots, D_m\}$ on D , the Gini index of D given this partitioning is

$$\text{Gini}_P(D) = \frac{|D_1|}{|D|} \text{Gini}(D_1) + \dots + \frac{|D_m|}{|D|} \text{Gini}(D_m)$$

- The reduction in impurity that would be incurred by the binary split is

$$\Delta\text{Gini}_P = \text{Gini}(D) - \text{Gini}_P(D)$$

Variance

- Variance is the expectation of the squared deviation of a random variable from its mean.
 - is a simple error measure for binary classification (i.e., two class labels, often represented by 0 and 1)
 - Given D , a data partition or a set of training tuples:

$$\text{Var}(D) = p(1 - p)$$

where p is the probability that a tuple in D belongs to class C_0 and is estimated by $|D_{C_0}|/|D|$.

Gain Ratio*

- Disadvantage of InfoGain: Tends to prefer splits that result in large number of partitions, each being small but pure.
- Recall that each split on node results in a partition $P = \{D_1, \dots, D_v\}$ on D , the set of records associated with this node.
- $\text{SplitInfo}(P) = - \sum_{i=1}^v \frac{|D_i|}{|D|} \log \left(\frac{|D_i|}{|D|} \right)$
- $\text{GainRatio} = \text{InfoGain}(P) / \text{SplitInfo}(P)$

Comparison of Impurity Measures

- All impurity measures return good results in general, but
 - **Information gain:**
 - biased towards multivalued attributes
 - **Gain ratio:**
 - tends to prefer unbalanced splits in which one partition is much smaller than the others
 - **Gini index:**
 - biased to multivalued attributes
 - has difficulty when the number of classes is large
 - tends to favor tests that result in equal-sized partitions and purity in both partitions
 - **Variance:**
 - suitable to binary classification, even though extension is possible

Advantages of Decision Tree Classifier

- Construction of the tree does not require any domain knowledge
- Can handle multidimensional data
- Representation of knowledge (as a decision tree) easy to assimilate by human
- The learning and classification steps are simple and fast
- Good accuracy in general.

Overfitting and Tree Pruning

- Overfitting: An induced tree may overfit the training data
 - Too many branches, some may reflect anomalies due to noise or outliers
 - Poor accuracy for unseen samples
- Two approaches to avoid overfitting
 - Pre-pruning: *Halt tree construction early*— do not split a node if this would result a measure falling below a threshold
 - Difficult to choose appropriate parameter thresholds
 - Post-pruning*: *Merge branches* from a “fully grown” tree—get a sequence of progressively pruned trees
 - Use a set of data *different* from the training data to decide which is the “best pruned tree”


Pre-pruning Sample Python Code (Multi-Split)

```
def chooseBestMultiSplit(dataSet, ops=(0.1,20)):
    tolG = ops[0]; tolN = ops[1]
    if (shape(dataSet)[0] < tolN):
        return None # exit
    numFeatures = len(dataSet[0]) - 1 # number of features
    baseEntropy = calcShannonEnt(dataSet)
    bestInfoGain = 0.0; bestFeature = -1
    for i in range(numFeatures): # iterate over all features
        uniqueVals = set([tuple[i] for tuple in dataSet])
        newEntropy = 0.0
        for value in uniqueVals:
            # "splitDataSet" function, implemented elsewhere, filters "dataset" such that the i-th
            # feature equals to "value"
            subDataSet = splitDataSet(dataSet, i, value)
            prob = len(subDataSet) / float(len(dataSet))
            newEntropy += prob * calcShannonEnt(subDataSet)
            infoGain = baseEntropy - newEntropy
            if (infoGain > bestInfoGain):
                bestInfoGain = infoGain; bestFeature = i
            if bestInfoGain < tolG:
                return None #exit
    return bestFeature # returns a feature index
```

ops is an optional argument. If the variance decrement is small than ops[0] or the size of the split dataset is small than ops[1], stop the split process. By default, ops=(0.5,4).

Random Forest: Model Ensemble for Decision Trees

Ensemble Methods:





- Ensemble methods
 - Use a combination of models
 - Combine a series of k learned models, M_1, M_2, \dots, M_k , with the aim of creating a combined model M^*
- Popular ensemble methods
 - Bagging: averaging the prediction over a collection of classifiers
 - Boosting: weighted vote with a collection of classifiers
 - Ensemble: combining a set of heterogeneous classifiers

Ensemble Methods:

- Advantages:
 - Increase accuracy: Miss classification occurs only when more than half of base classifiers predict incorrectly (even better if the base classifiers are less correlated).
 - Can deal with data in sheer volume (too many records or attributes)
 - Can run in parallel

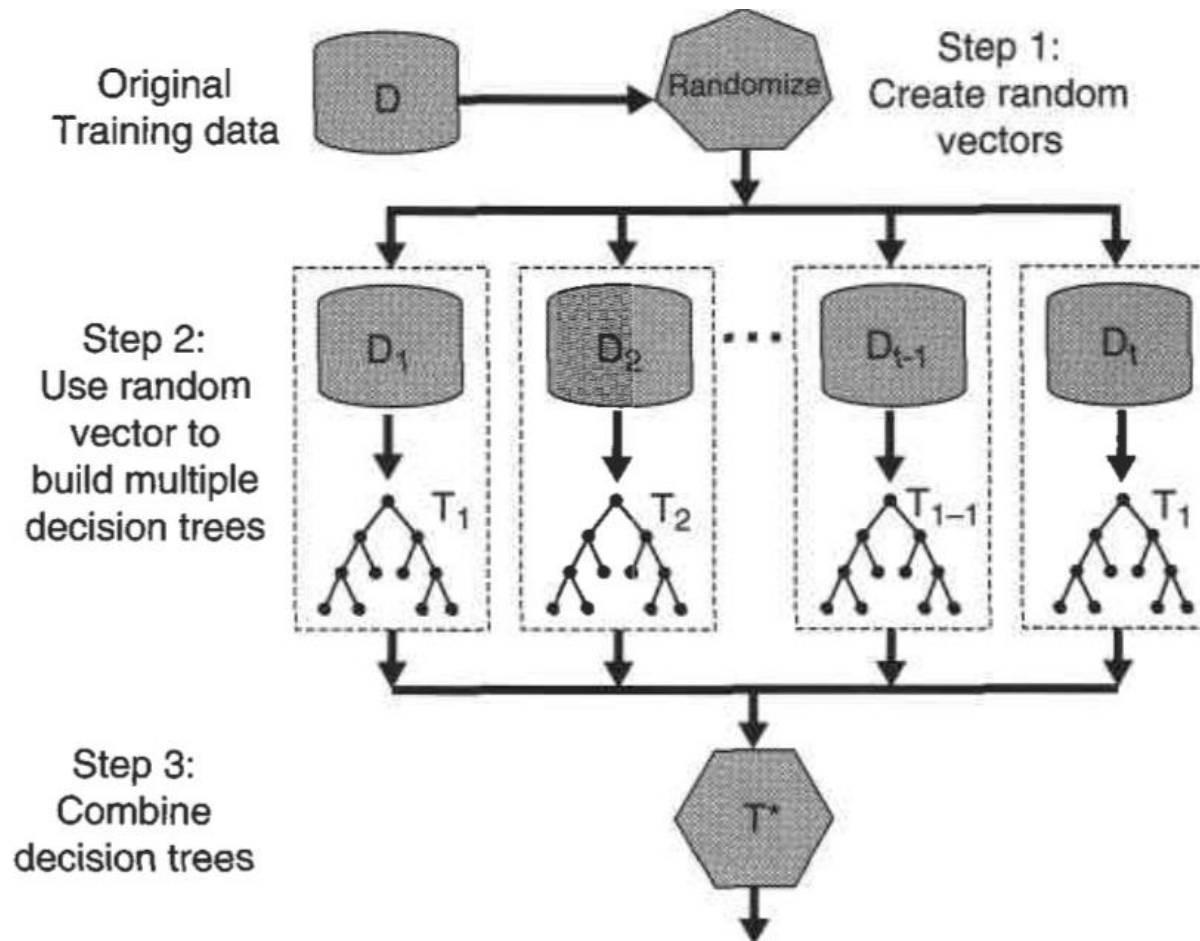
Decision boundary by (a) a single decision tree and (b) a random forest



Random Forest

- Random Forest is a class of ensemble methods specifically designed for decision tree classifiers.
 - It combines the predictions made by multiple decision trees.
 - Each tree is generated randomly based on the training tuples.
 - The final prediction output is produced by a voting function.
- A properly built random forest tends to be more accurate and less biased than individual decision tree classifiers.
 - The accuracy of RF depends on the *strength* of individual classifiers (trees) and a measure of *dependence* between them.
- But the computational cost grows as the number of trees in the forest increases.

Random Forest



Random Forest

- There are 3 common ways to associate randomization with decision trees.
- **(1) Bagging:** Given a set D of d tuples, bagging works as follows. For iteration i ($i = 1, 2, \dots, k$), a training set D_i of d tuples is sampled *with replacement* from the original set D .
- Note that some of the original tuples of D may not be included in D_i , whereas others may occur more than once.
- A decision tree M_i is learned for each training set, D_i . To classify an unknown tuple X , each classifier M_i returns its class prediction, which counts as one vote.
- The bagged classifier, say, M_* counts the votes and assigns the class with the most votes to X .
- Random Forests can handle datasets that don't fit in memory

Random Forest

- **(2) Forest-RI** (*random input selection*)
 - When building the tree, randomly select F attributes (features) that are used to determine the split at each node, where F is much smaller than the number of available attributes.
 - Useful when the number of attributes is large

Random Forest

- **(3) Forest-RC** (*random linear combinations*)
 - creates new attributes (features) that are a linear combination of the existing attributes.
 - that is, randomly selected and added together with coefficients that are uniform random numbers on $[-1,1]$.
 - Useful when the number of attributes is small or large

Summary

- Decision Tree Classifier
 - Theory
 - Implementation
 - Tree Pruning
- Random Forest