

U

O

W

Data Processing with Apache Spark

– Handling Massive Data Set (Part I)

CSCI316: Big Data Mining Techniques and Implementation



UNIVERSITY
OF WOLLONGONG
AUSTRALIA

Contents

MapReduce and Spark Model

Using Spark

MapReduce and Spark Model

Processing Massive Data

- Processing massive datasets requires long runtime
 - One solution to speed up the computation is *parallelism* (or *distributed computing*)
- **MapReduce:** One preeminent model of parallel computation
- A very simple model:
 - One *Map* stage, which performs simple mapping-alike operations to produce key-value pairs
 - One intermediate stage, which merges key-value pairs per key
 - One *Reduce* stage, which performs aggregation-alike operations per key
- However...
 - from the algorithm point of view: very powerful!
 - from the implementation point of view: very suitable for a computing cluster

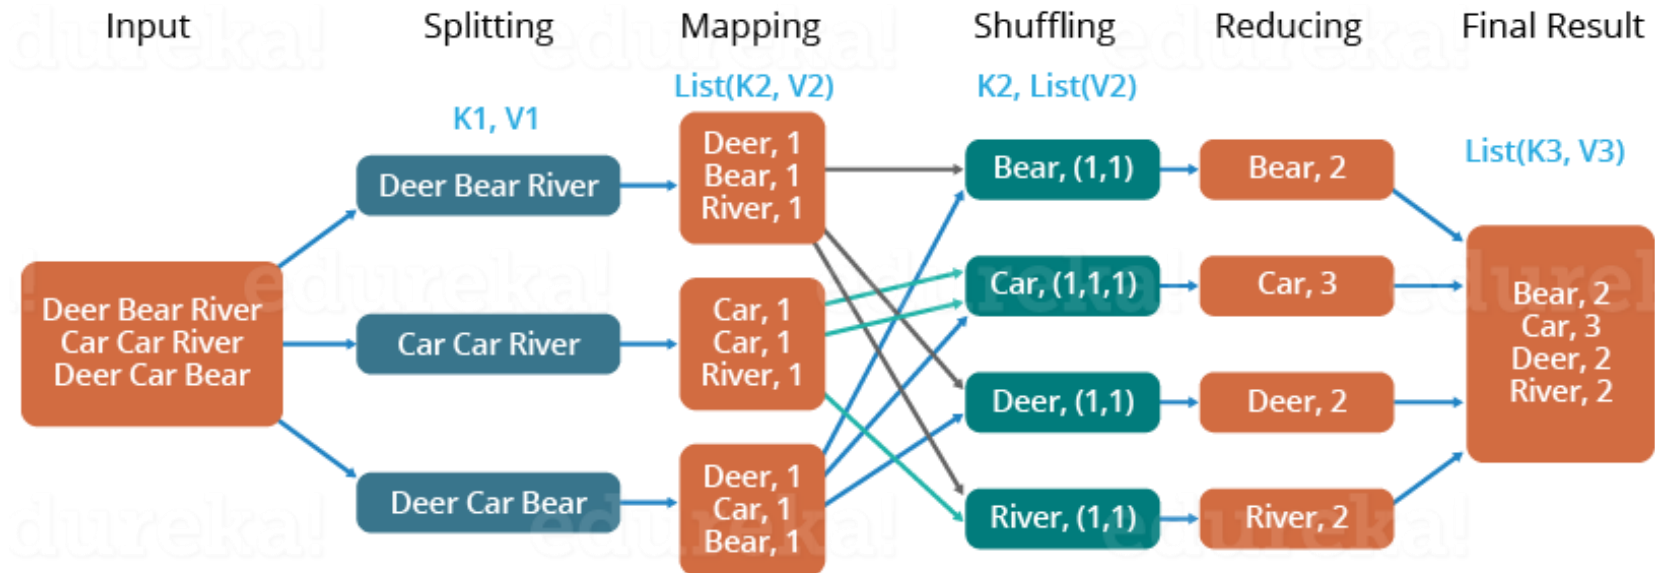


MapReduce Workflow: The WordCount Example

- MapReduce uses (key, value) as the basic data structure.

The Overall MapReduce Word Count Process

edureka!




Relational-Algebra Operations in MapReduce

- Although big data frameworks are not traditionally DB systems, **relational-algebra operations** are useful, especially in preprocessing.
- Recall that a *relation* is a table. We call the column headers as *attributes* and the rows as *tuples*. The bag of attributes of a relation is called its *schema*.
- We use $R[A_1, \dots, A_n]$ to denote a relation R with schema A_1, \dots, A_n .
- **Selections:** Apply a condition C to each tuple in the relation and produce as output only the tuples that satisfy C .
- Computing selections in MapReduce
 - The Map function: For each tuple t in R , test if it satisfies C . If so the mapper produce the key-value pair (t, t) ; otherwise, it produces nothing.
 - The Reduce function: The identity function.

Relational-Algebra Operations in MapReduce

- **Natural Join**


- Given two relations, compare each pair of tuples, one from each relation. If the two tuples agree on all the attributes that are common to the two schemas, then produce a tuple that has components for each of the attributes in either schema or both.



- Computing Natural Join in MapReduce

- **Map function.** For each tuple (a, b) in $R[A, B]$, produce the key-value pair $b: (R, a)$. For each tuple (b, c) in $S[B, C]$, produces the key-value pair $b: (S, c)$.
- **Reduce function.** Each key value b will be associated with a list of pairs that are of the form (R, a) or (S, c) . **Construct all tuples (a, b, c) if both (R, a) and (S, c) appear in the list that b is associated with.**

Example



- Map function

- From table R.

$(a_1, b_1) \mapsto b_1: (R, a_1)$

$(a_2, b_1) \mapsto b_1: (R, a_2)$

$(a_3, b_2) \mapsto b_2: (R, a_3)$

$(a_4, b_3) \mapsto b_3: (R, a_4)$

- From table S.

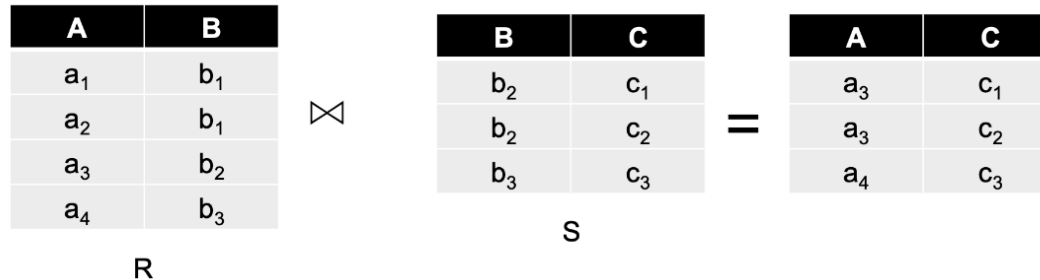
$(b_2, c_1) \mapsto b_2: (S, c_1)$

$(b_2, c_2) \mapsto b_2: (S, c_2)$

$(b_3, c_3) \mapsto b_3: (S, c_3)$

Example

- Reduce function
 - $b_1 : (R, a_1), (S, a_2) \mapsto \text{None}$
 - $b_2 : (R, a_3), (S, c_1), (S, c_2) \mapsto (a_3, c_1), (a_3, c_2)$
 - $b_3 : (R, a_4), (S, c_3) \mapsto (a_4, c_3)$



Relational-Algebra Operations in MapReduce

- You can verify that other all other usual relational-algebra operations, such as projection, union, intersection, grouping and left/right/outer-joins, can be expressed in the model of MapReduce
- **Conclusion:** *All common SQL queries can be implemented with MapReduce*

Reference: Sec. 2.3.3 – 2.3.8, J. Leskovec, A. Rajaraman, J. Ullman. Mining of Massive Datasets, 3rd Edition

Matrix-Matrix Multiplication in MapReduce

- If \mathbf{M} is a matrix with element m_{ij} in row i and column j , and \mathbf{N} is a matrix with element n_{jk} in row j and column k , then the product $\mathbf{P} = \mathbf{MN}$ is the matrix with element p_{ik} in row i and column k , where $p_{ik} = \sum_j m_{ij}n_{jk}$
 - Note that the number of columns of \mathbf{M} must equals to the number of row in \mathbf{N} .
- We can view as a matrix as a *relation* with three attributes: the row number, the column number, and the value in that row and column.
- Thus, $\mathbf{M}[I, J, V]$ with tuples (i, j, m_{ij}) and $\mathbf{N}[J, K, W]$ with tuples (j, k, n_{jk})
- Adopting this idea, can develop a *two-stage* MapReduce job for Matrix-Matrix Multiplication.

Matrix-Matrix Multiplication in MapReduce

The first pair of MapReduce functions:

- **Map Function A:** For each matrix element m_{ij} , produce the key value pair $j : (M, i, m_{ij})$. Likewise, for each matrix element n_{jk} , produce the key value pair $j : (N, k, n_{jk})$.
- **Reduce Function A:** For each key j , examine its list of associated values. For each value that comes from M, say (M, i, m_{ij}) , and each value that comes from N, say (N, k, n_{jk}) , produce a key-value pair with **key equal to (i, k)** and value equal to the product of these elements, $m_{ij}n_{jk}$.

Matrix-Matrix Multiplication in MapReduce

The second MapReduce performs a grouping and aggregation applied to the output of the first MapReduce.

- The **Map Function B**: This function is just the identity. That is, for every input element with key (i, k) and value v , produce exactly this key-value pair.
- The **Reduce Function B**: For each key (i, k) , produce the sum of the list of values associated with this key. The result is a pair $(i, k) : v$, where v is the value of the element in row i and column k of the matrix $\mathbf{P} = \mathbf{MN}$.


Reference: Sec. 2.3.9, J. Leskovec, A. Rajaraman, J. Ullman. Mining of Massive Datasets, 3rd Edition

The kNN Classifier in MapReduce

- How to write a kNN classifier in MapReduce?
- Recall the movie example in the second lecture.
- The basic idea:
 - Mapper returns **<key1, val1>** where **key1** is a movie name and **val1** is the distance to the unknown movie
 - Reducer returns **<key2, val2>** where **key2** is **null** (not important) and **val2** is a list of k nearest movies (to the unknown movie) and the distances
 - ❖ Can use a combiner to improve the performance (why?)
 - Finally, a voting function is used based on **val2** to determine the class for the unknown movie.

MapReduce to DAG

- A **directed acyclic graph** (DAG) extends MapReduce
 - from a simple two-step model (with a mapper and a reducer) to a orchestration of any steps that form a DAG
 - Although in theory is possible to pipeline MapReduce jobs to form any workflow, however...
 - You need to store the temporal output of intermediate jobs in HDFS (which is a natural idea in MapReduce) rather than keep it in memory
 - Workflow system in Apache Spark:




Using Spark

Distributed vs. single-machine computing

Spreadsheet on
a single machine



Table or Data Frame
partitioned across servers
in a data center




Spark's Abstractions and APIs

- DataFrame—high-level *structured* abstraction
 - Similar to Pandas Dataframe but support distributed computing
 - Intuitively, a DataFrame is a table of data with rows and columns
 - which may be stored in a *single or multiple* machines
 - There is a scheme that defines the meta information (e.g. data types) for the columns
- *Resilient Distributed Dataset (RDD)—low-level abstraction
 - More control, sometimes more flexible, but less efficient than DataFrame
 - Can convert to a DataFrame.
- Immutability: Both DataFrame and RDD are immutable.
 - Instead of altering elements in a DataFrame or RDD, you create a new one

Transformation and Action

- End users operate on DataFrame as if the data is on a single computer
- Two kinds of operations: Transformations and Actions
 - Transformations create another DataFrame/RDD (e.g., create a new row)
 - Actions produce a computational result (e.g., count the rows)
 - A spark application can be viewed as a DAG (direct acyclic graph) of transformations and actions

Spark Data Analytics Pipeline as a DAG



Lazy evaluation: Spark don't evaluate the DataFrame/RDD until it has to, e.g., an action is performed.

Spark's DataFrame API (PySpark)

- See supplementary materials for some examples.
- Comprehensive API reference:
<https://spark.apache.org/docs/latest/api/python/reference/pyspark.sql/dataframe.html#>

Summary

- MapReduce Model
 - A powerful computation model for processing massive data
 - Hadoop's MapReduce Framework
- Spark's DAG Model
 - A DAG model consisting of a series of transformations and an action
 - Spark's rich set of APIs