

Artificial Neural Networks and TensorFlow - Comprehensive Exam Notes

1. TensorFlow Overview

What is TensorFlow?

- **Definition:** Python-friendly open source library for numerical computation
- **Purpose:** Well-suited for large-scale machine learning and deep learning
- **Key Features:**
 - Define computation graphs in Python
 - Breaks graphs into chunks for parallel execution
 - Supports multiple CPU, GPU, and TPU processing

TensorFlow with Keras

- Keras is the high-level API for TensorFlow
 - Provides user-friendly interface for building neural networks
 - Simplifies model creation, training, and evaluation
-

2. Linear Threshold Unit (LTU)

Basic Structure

- **Inputs:** Numerical values (x_1, x_2, \dots, x_n)
- **Weights:** Each input has an associated weight (w_1, w_2, \dots, w_n)
- **Computation:**
 - Weighted sum: $z = w_1x_1 + w_2x_2 + \dots + w_nx_n = \mathbf{w}^T \mathbf{x}$
 - Output: $\hat{y} = \text{step}(z) = \text{step}(\mathbf{w}^T \mathbf{x})$
- **Step Function:** $\text{step}(z) = 0$ if $z < 0$, otherwise 1

Mathematical Representation

- Vector form: $z = \mathbf{w}^T \mathbf{x}$
 - Decision boundary created by step function
 - Forms basis for more complex neural networks
-

3. Perceptron

Architecture

- **Structure:** Single layer of neurons
- **Bias:** Added as fixed input value of 1
- **Neuron j computation:** $z_j = w_{j1}x_1 + \dots + w_{jn}x_n + b = w_j^T x + b$

Activation Functions (replacing step function)

1. **ReLU (Rectified Linear Unit):** $\text{ReLU}(z_j) = \max(0, z_j)$
2. **Sigmoid (Logistic):** $\sigma(z_j) = e^{z_j} / (1 + e^{z_j}) = 1 / (1 + e^{-z_j})$
3. **Hyperbolic Tangent:** $\tanh(z_j) = 2\sigma(2z_j) - 1$

Training Process

- **Weight Update:** $w_{ij}(\text{new}) = w_{ij} - \alpha \cdot \partial J(w_j, b_j) / \partial w_{ij}$
 - **Bias Update:** $b_j(\text{new}) = b_j - \alpha \cdot \partial J(w_j, b_j) / \partial b_j$
 - **Learning Rate:** α (step size parameter)
-

4. Multi-Layer Perceptron (MLP)

Architecture Components

- **Input Layer:** Conceptual layer that forwards inputs
- **Hidden Layer(s):** One or more layers between input and output
- **Output Layer:** Produces final predictions
- **Connections:** Fully connected layers (except output layer includes bias)

Key Properties

- **Universal Approximation:** Sufficiently large MLP can approximate any continuous function
- **Feedforward:** Information flows in one direction (input \rightarrow hidden \rightarrow output)
- **Bias Neurons:** Special neurons that always output 1

Hidden Layer Computation

- **Input:** $h = (h_1, \dots, h_n)$ from previous layer
- **Computation:** $z_j = w_j^T h + b$
- **Activation Functions:** Sigmoid, Tanh, or ReLU

5. Output Layer Configurations

For Regression

- **Linear Output:** $\hat{y} = w_j^T h + b_j$ (no activation)
- **Bounded Output:** $\hat{y} = \sigma(w_j^T h + b_j)$ (sigmoid activation)
- **Positive Output:** $\hat{y} = \text{ReLU}(w_j^T h + b_j)$

For Classification

- **Softmax Function:**
 - Intermediate: $z_j = w_j^T h + b_j$
 - Output: $\hat{y}_j = \text{softmax}_j(z) = e^{z_j} / \sum_{i=1}^n e^{z_i}$
 - Interpretation: \hat{y}_j represents probability of class j
-

6. Cost Functions

Regression

- **Mean Squared Error (MSE):** Standard loss function for regression problems
- **Formula:** $J = (1/m) \sum_{i=1}^m (y_i - \hat{y}_i)^2$

Classification

- **Cross-Entropy Loss:**
 - Single sample: $\text{cross_entropy}(y_i, \hat{y}_i) = -\sum_j y_{ji} \log(\hat{y}_{ji})$
 - Total cost: $\text{cost}(y, \hat{y}) = (1/m) \sum_{i=1}^m \text{cross_entropy}(y_i, \hat{y}_i)$
 - **Binary Classification:** Binary cross-entropy with sigmoid activation
-

7. Training Process (Backpropagation)

Forward Pass

1. **Input Processing:** Feed training instance x to network
2. **Layer-by-layer Computation:** Calculate outputs for each layer
3. **Prediction:** Compute final output $\hat{y} = f(x)$
4. **Error Calculation:** Measure $\text{cost}(y, \hat{y})$

Backward Pass

1. **Output Layer Error:** Calculate error contribution from each output neuron
2. **Hidden Layer Error:** Work backwards to measure error contribution from each hidden neuron
3. **Gradient Calculation:** Compute gradients for all weights and biases
4. **Weight Update:** Apply gradient descent to reduce error

Key Concepts

- **Chain Rule:** Used to compute gradients through multiple layers
 - **Gradient Descent:** Iterative optimization algorithm
 - **Learning Rate:** Controls step size in weight updates
-

8. Keras Implementation

Dataset Preparation

```
python

# Example: Fashion MNIST
fashion_mnist = keras.datasets.fashion_mnist
(X_train_full, y_train_full), (X_test, y_test) = fashion_mnist.load_data()

# Scaling and splitting
X_valid, X_train = X_train_full[:5000] / 255.0, X_train_full[5000:] / 255.0
y_valid, y_train = y_train_full[:5000], y_train_full[5000:]
```

Model Architecture (Sequential API)

```
python

model = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[28, 28]), # 2D to 1D conversion
    keras.layers.Dense(300, activation="relu"), # Hidden layer 1
    keras.layers.Dense(100, activation="relu"), # Hidden layer 2
    keras.layers.Dense(10, activation="softmax") # Output layer
])
```

Model Compilation

```
python
```

```
model.compile(  
    loss="sparse_categorical_crossentropy", # For sparse labels  
    optimizer="sgd",                      # Stochastic Gradient Descent  
    metrics=["accuracy"]                  # Performance metric  
)
```

Training

```
python
```

```
history = model.fit(  
    X_train, y_train,  
    epochs=30,  
    validation_data=(X_valid, y_valid)  
)
```

Evaluation and Prediction

```
python
```

```
# Evaluate model  
model.evaluate(X_test, y_test)  
  
# Make predictions  
y_proba = model.predict(X_new)  
y_pred = np.argmax(y_proba, axis=-1)
```

9. Hyperparameter Configurations

Regression Networks

Parameter	Typical Values
Input neurons	One per feature
Hidden layers	1-20 layers
Neurons per layer	10x-100x input features
Output neurons	1 per target variable
Hidden activation	ReLU, Sigmoid, Tanh
Output activation	None, ReLU, Sigmoid/Tanh
Loss function	MSE

Classification Networks

Task	Binary	Multi-label	Multi-class
Output neurons	1	1 per label	1 per class
Output activation	Sigmoid	Sigmoid	Softmax
Loss function	Binary Cross-entropy	Binary Cross-entropy	Categorical Cross-entropy

10. Advanced Techniques

Regularization

- **L1 Regularization:** $\tilde{J}(W) = J(W) + \lambda_1 \sum_i |w_i|$
- **L2 Regularization:** $\tilde{J}(W) = J(W) + \lambda_2 \sum_i w_i^2$
- **Combined:** $\tilde{J}(W) = J(W) + \lambda_1 \sum_i |w_i| + \lambda_2 \sum_i w_i^2$

Implementation in Keras

```
python
keras.layers.Dense(
    4,
    activation="relu",
    kernel_regularizer=keras.regularizers.l1_l2(l1=0.01, l2=0.01)
)
```

Early Stopping

- **Purpose:** Prevent overfitting
- **Method:** Stop training when validation error stops decreasing

- **Implementation:** Monitor validation loss during training
-

11. Hyperparameter Tuning

Grid Search with Keras Tuner

python

```
def build_model(hp):
    n_hidden = hp.Int("n_hidden", min_value=0, max_value=8, default=2)
    n_neurons = hp.Int("n_neurons", min_value=16, max_value=256)
    learning_rate = hp.Float("learning_rate", min_value=1e-4, max_value=1e-2)

    model = tf.keras.Sequential()
    model.add(tf.keras.layers.Flatten())
    for _ in range(n_hidden):
        model.add(tf.keras.layers.Dense(n_neurons, activation="relu"))
    model.add(tf.keras.layers.Dense(10, activation="softmax"))

    model.compile(
        loss="sparse_categorical_crossentropy",
        optimizer=tf.keras.optimizers.SGD(learning_rate=learning_rate),
        metrics=["accuracy"]
    )
    return model
```

12. Loss Functions in Keras

Common Loss Functions

- **sparse_categorical_crossentropy:** For integer labels (0, 1, 2, ...)
- **categorical_crossentropy:** For one-hot encoded labels
- **binary_crossentropy:** For binary classification
- **mse:** Mean Squared Error for regression

Selection Criteria

- **Sparse labels:** Use sparse_categorical_crossentropy
- **One-hot labels:** Use categorical_crossentropy
- **Binary tasks:** Use binary_crossentropy
- **Regression:** Use mse or other regression losses

13. Optimizers

Available Optimizers in Keras

- **SGD**: Standard Stochastic Gradient Descent
- **Adam**: Adaptive Moment Estimation
- **RMSprop**: Root Mean Square Propagation
- **Adagrad**: Adaptive Gradient Algorithm
- **Adadelta**: Extension of Adagrad

Choosing Optimizers

- **SGD**: Simple, interpretable, good baseline
 - **Adam**: Generally good performance, adaptive learning rates
 - **RMSprop**: Good for recurrent neural networks
-

14. Activation Functions

Available Functions

- **ReLU**: Most common for hidden layers
- **Sigmoid**: Output layer for binary classification
- **Softmax**: Output layer for multi-class classification
- **Tanh**: Alternative to sigmoid
- **Linear**: No activation (regression output)

Selection Guidelines

- **Hidden layers**: Usually ReLU
 - **Binary output**: Sigmoid
 - **Multi-class output**: Softmax
 - **Regression output**: Linear or ReLU (for positive values)
-

15. Initialization Strategies

Available Initializers

- **GlorotNormal/GlorotUniform**: Good for sigmoid/tanh

- **HeNormal**: Good for ReLU
- **Constant**: Initialize to specific value
- **Random**: Various random initializations

Usage

python

```
keras.layers.Dense(10, activation="relu", kernel_initializer="he_normal")
```

16. Key Exam Tips

Important Concepts to Remember

1. **Architecture Design**: Know when to use different layer types and sizes
2. **Activation Functions**: Understand when to use each type
3. **Loss Functions**: Match loss function to problem type
4. **Regularization**: Understand L1/L2 regularization effects
5. **Training Process**: Understand forward/backward pass
6. **Hyperparameter Tuning**: Know common ranges and selection criteria

Common Exam Questions

1. **Design networks**: Given a problem, specify architecture
2. **Choose parameters**: Select appropriate loss, optimizer, activation
3. **Debug training**: Identify overfitting/underfitting issues
4. **Mathematical understanding**: Compute forward/backward pass
5. **Implementation**: Write Keras code for specific architectures

Problem-Solving Approach

1. **Identify problem type**: Regression vs. classification
2. **Design architecture**: Input → Hidden → Output layers
3. **Choose components**: Activation, loss, optimizer
4. **Consider regularization**: If overfitting is likely
5. **Plan evaluation**: Metrics and validation strategy